Politecnico di Torino

Master of Science in Computer Engineering

Tesi di Laurea Magistrale

Rethinking Automotive Software Development: Exploring Software Defined Vehicle and its potential

Supervisors
prof. Danilo Bazzanella
dott.sa Piera Limonet

Candidate
Lorenzo Sciara

Anno Accademico 2023-2024

Summary

This thesis project delves into the analysis of contemporary connected vehicle platforms, focusing on the benefits and challenges associated with these advanced solutions and emphasising aspects of safety and flexibility. A key trend in the current automotive sector is the prospect of transforming the car from a hardware-focused product to a software-driven device. The technology of choice for leading software development and production companies driving this change is the Software Defined Vehicle (SDV).

The primary objective of the thesis is to apply this paradigm to the development of a simulator for a vehicle control unit responsible for collecting telemetric data from the vehicle. The implementation of the simulator involves an in-depth analysis of the drawbacks of the automotive software production industry and the advantages of the Software Defined Vehicle solution. The simulator implementation also includes the creation of a scaled-down version of a connected vehicle platform, storage infrastructure and example application.

Using the Amazon Web Services (AWS), an environment in the cloud is established for the development of the necessary software for the operation of the vehicle control unit. Development of the vehicle control unit simulator is carried out, including client connectivity to interact with the cloud platform, telemetry generation, logic for remote operations, and optional applications. The final phase involves testing the simulator on compatible hardware to validate its functionality and performance.

The successful completion of this project in collaboration with Storm Reply, not only highlights the potential of the software-defined vehicle paradigm as a leading force in the future of the automotive sector, but also explores the economic, safety and security benefits associated with its adoption, paving the way for significant progress in the field and ensuring an advanced and safe end-user experience.

Acknowledgements

Acknowledgement (optional)

Contents

List of Figures

List of Tables

Listings

Chapter 1

Introduction

1.1 Context

The automotive industry stands out as one of the fastest-growing sectors, playing a significant role as both an employer and an investor in research and development; at the same time, it represents one of the most crucial domains for the European Union's economy. As reported in the article [**?**], in 2015, 21 million motor vehicles of all types were produced in Europe, representing a 23% share in the global production of more than 90 million units.

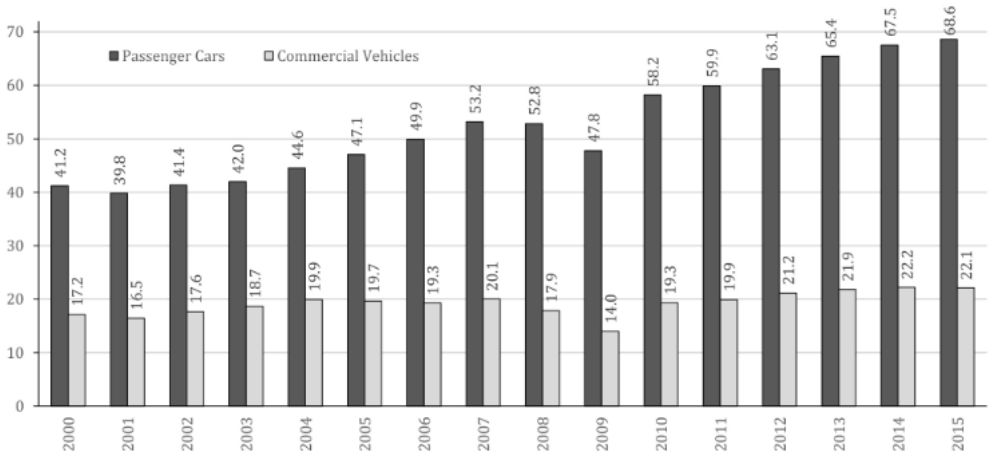

Figure 1.1. World automobile production in million vehicles [**?**]

In the evolving landscape of automotive technology, the imperative for automotive companies extends beyond the traditional realms of mechanical engineering to encompass a crucial reliance on both software and hardware components for vehicle construction. A glimpse into the intricate web of modern cars, as illustrated in Figure **??**, reveals a mosaic of hundreds of distinct processors interfacing at various levels, earning contemporary vehicles the moniker of "Computers on wheels."

However, the proliferation of processors within vehicles, orchestrating communication to manage diverse components, presents a formidable challenge; each component often integrates a processor with unique logics, diverging from the logics

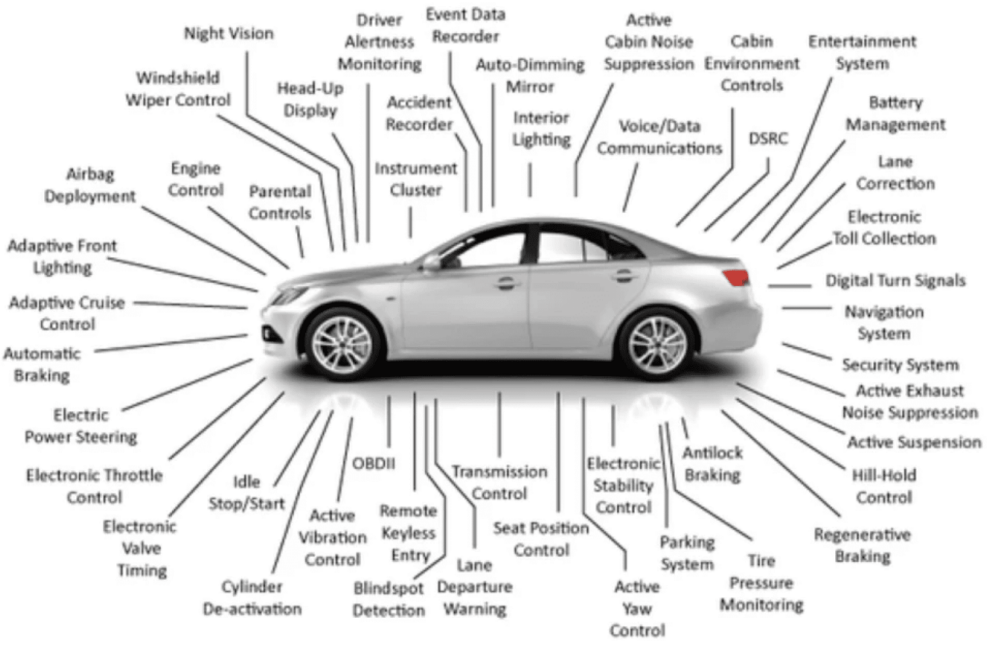

Figure 1.2. An incomplete overview of computers in a modern car [**?**]

embedded in processors of other components. Complicating matters further, these components are frequently supplied by companies with proprietary management logics, not readily accessible to the automotive companies themselves.

In addressing this intricate scenario, the transformative concept of a Software Defined Vehicle (SDV) comes to the forefront. Defined as "any vehicle that manages its operations, adds functionality, and enables new features primarily or entirely through software" [**?**], the notion of SDV offers a comprehensive solution to the challenges posed by the intricate interplay of software and hardware in modern vehicles.

Effectively navigating the development of SDV technology necessitates a collaborative approach across diverse companies, particularly in the realms of hardware and cloud computing. This collaborative synergy is exemplified in the realization of our project, made possible through the partnership with Storm Reply.

1.2 Company

Leveraging extensive experience in the cloud industry and fostering deep-rooted relationships within the automotive sector, Storm Reply stands out as the ideal choice to lead the project discussed in this thesis. A key player in the Reply group, Storm Reply specializes in designing and implementing innovative Cloud-based solutions and services [**?**].

With a diverse clientele spanning various sectors, notably the automotive industry, the company's expertise played a pivotal role in comprehensively understanding the project's context and internal dynamics. This profound knowledge served as the cornerstone for developing a tangible exemplification of the infrastructure.

Figure 1.3. Logo of the partenr company of the project

A point of pride for Storm Reply is its recognition as an Amazon Web Services (AWS) Premier Consulting Partner since 2014, ranking among the top Amazon Partners globally. This distinctive characteristic underscores the decision to develop the infrastructure using Amazon Web Services.

According to the official AWS description page [**?**] the AWS Cloud spans 102 Availability Zones within 32 geographic Regions around the world and servs 245 countries and territories. With millions of active customers and tens of thousands of partners globally, AWS has the largest and most dynamic ecosystem. AWS is evaluated as a Leader in the 2022 Gartner Magic Quadrant for Cloud Infrastructure and Platform Services, placed highest in Ability to Execute axis of measurement among the top 8 vendors named in the report.

Figure 1.4. Here are a series of market research reports published by IT consulting firm Gartner that rely on proprietary qualitative data analysis methods to demonstrate market trends, such as direction, maturity and participants. [**?**]

The infrastructure exhibits several key attributes contributing to its robustness and efficiency:

- **Security:** The infrastructure undergoes 24/7 monitoring to ensure the confidentiality, integrity, and availability of data. All data flowing across the AWS global network is automatically encrypted at the physical layer before leaving secured facilities.
- **Availability:** To ensure high availability and isolate potential issues, applications can be partitioned across multiple AZs (Availability Zones) within the same region, creating fully isolated infrastructure partitions.
- **Performance:** AWS Regions offer low latency, low packet loss, and high overall network quality. This is achieved through a fully redundant 100 GbE fiber network backbone, often providing terabits of capacity between Regions.

- **Scalability:** The AWS Global Infrastructure allows companies to take advantage of the virtually infinite scalability of the cloud. This enables customers to provision resources based on actual needs, with the ability to instantly scale up or down according to business requirements.
- **Flexibility:** The AWS Global Infrastructure provides flexibility in choosing where and how workloads are run, whether globally, with single-digit millisecond latencies, or on-premises.
- **Global Footprint:** AWS boasts the largest global infrastructure footprint, continually expanding at a significant rate.

1.3 Thesis Goal

In the automotive context, the use of Software Defined Vehicle (SDV) plays a crucial role in terms of costs, innovation, and safety. The goal of the thesis intertwine with the opportunities provided by Software Defined Vehicle technology, addressing the primary challenge of managing the current difficulties associated with the presence of various specialized hardware platforms on the same vehicle.

The central objective of this thesis is to propose a Software Defined Vehicle solution capable of eliminating various phases of the software production pipeline. This would result in significant time and cost savings, enabling the investment of these resources in other sectors. Since, by definition, a Software Defined Vehicle is characterized by the ability to undergo software updates dynamically and flexibly, this solution offers significant security advantages in various aspects:

1. Human Safety Critical Security: From the moment that a vehicle can be classified as safety critical (as it is reported in the standard ISO 26262-1:2018 of the ISO society where is said that "safety is one of the key issues in the development of road vehicles" [**?**]), the elimination of software vulnerabilities related to the vehicle's systems is crucial for the overall safety of the vehicle itself.
2. Intrinsic Software Security: This approach allows for the prevention and resolution of vulnerabilities unknown at the time of software design, contributing to ensuring a high standard of security.

Consequently, the use of Software Defined Vehicle aims to completely separate software and hardware, allowing the production of high-level software on entirely generalized hardware systems. This results in significant savings in terms of time and money for hardware production, along with providing an advantage in terms of security due to the simplification of software.

For example, as demonstrated by NIST in the research on the Analysis Of The Impact Of Software Complexity [**?**], the increase in software complexity in different cases results in less analyzable programs. In some instances, the same vulnerability analysis tool may detect vulnerabilities, while in others, analyzing the same code, it may not.

From a practical standpoint, the project's goal is to provide, through the use of AWS services, a cloud infrastructure capable of managing the Software Defined Vehicle both in terms of software production and data analysis.

Chapter 2

State-of-the-Art Analysis

The following chapter constitutes an in-depth exploration of current technologies and methodologies within the automotive industry, with a specific focus on the complexity of vehicular software development. Firstly, the current automotive landscape will be examined, providing a detailed insight into challenges associated with software development in vehicles.

Subsequently, through meticulous analysis of scientific publications, technical reports, and practical implementations, the chapter delves into the radical transformation of the automotive sector facilitated by the concept of Software Defined Vehicle (SDV). This technology, crucial for technological progress and vehicular safety, will be explored from various perspectives. Particularly, the synergy between Cloud, software, and hardware will be investigated, highlighting solutions proposed by major industry players and analyzing their applications, benefits, and limitations.

The objective is to offer a comprehensive overview of current dynamics, emphasizing the pivotal role of SDV in the evolution of the automotive industry.

2.1 Current Automotive Software Development

In the past, the automotive industry advanced primarily through the development of technologies in mechanical engineering, focusing on perfecting combustion engines. Nowadays, the paradigm has radically changed due to multiple factors, including electrification, automation, shared mobility, and connected mobility.

Software technology development in the automotive field can be metaphorically compared to what has happened in smartphone development, as highlighted in the manifesto document regarding Bosch's Software Defined Vehicle (SDV) [**?**].

The ultimate goal is to achieve simple and user-friendly devices that fully meet the user's needs. Currently, many customers express dissatisfaction because their cars do not offer the same functionality and ease of use common in smartphones. Many ask: Why can't my $50,000 car perform the same tasks as my $300 smartphone?

A key difference between the automotive and smartphone industries is the level of complexity, which brings with it a number of issues.

2.1.1 difficulties

We can analyse in depth the problems of the current automotive software that is being developed via 4 main difficulties:

- **Specialized Hardware:** Today's vehicles are still complex systems of systems. Each subsystem in a car, from brakes to transmission, is a complex entity, supplied by a different manufacturer and integrated with a unique software architecture. The level of complexity and the need for seamless interoperability between systems far exceeds that of today's smartphones.
- **Time:** The software production pipeline involves many development and testing steps with a not inconsiderable amount of time spent on each one. This is greatly increased by the presence of different components, so development time must be considered for each different unit of the system.
- **Cost:** The complexity of the software systems in vehicles entails very high costs, aggravated by the fact that the test phase is often carried out directly on the boards (for hardware requirements), which means a much longer production process, especially in the event of errors.

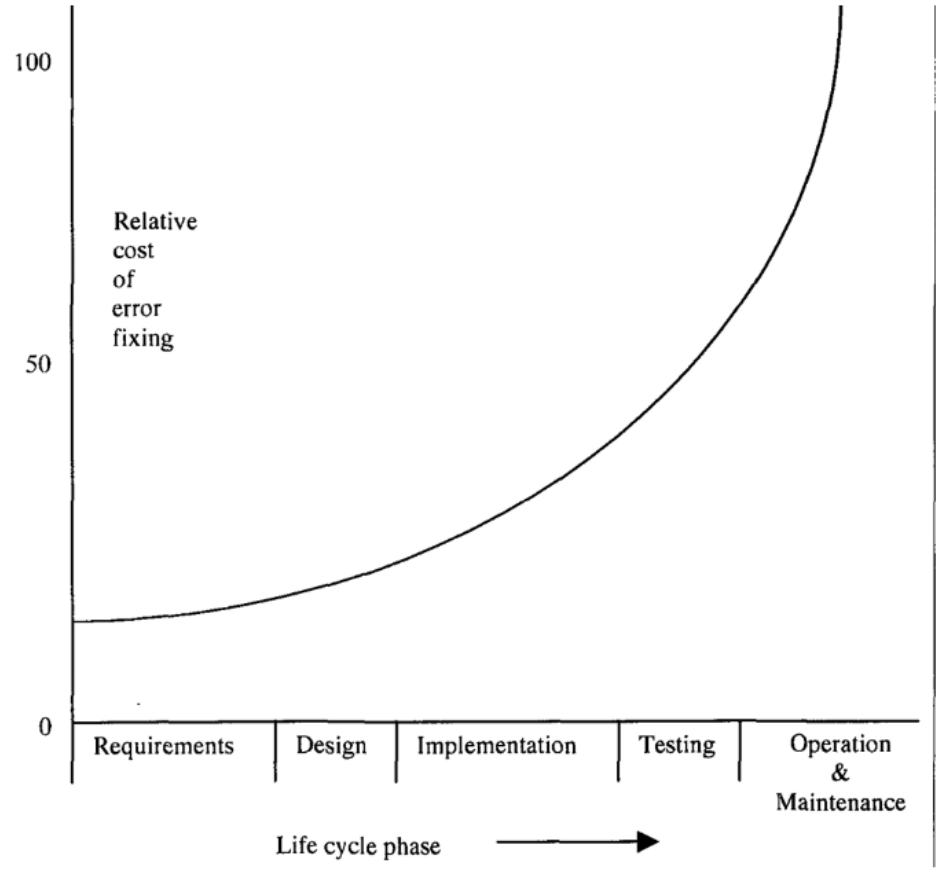

Figure 2.1. Cost of fixing errors increases in later phases of the life cycle [**?**]

- **Human Safety Security:** Automotive embedded software must meet stringent reliability and security requirements, while delivering performance and a reasonable memory footprint. To develop automotive embedded software, you need the right tools that meet safety and security standards to evaluate, prototype and test your software.

What lessons can be drawn from the study of barriers that can be applied to the vehicle lifecycle? Historically, the vehicle lifecycle has been characterised by the simultaneous production and deployment of tightly integrated hardware and software. Once the vehicle was in the hands of the consumer, its characteristics remained largely unchanged until the end of its life. However, the SDV paradigm introduces the possibility of decoupling hardware and software release dates, a prerequisite for adopting a digital-first approach. This approach brings the design and virtual validation of the digital vehicle experience to the forefront of the lifecycle. It also requires the application of the digital-first concept, which means that new ideas for the vehicle experience are first explored in virtual environments to ensure early user feedback, long before any custom hardware needs to be developed or a physical test vehicle is available. Digital first is the application of design thinking and lean startup principles, originally rooted in internet culture, to the tangible realm of automotive development.

2.2 Introduction to Software Defined Vehicle

The Software Defined Vehicle represents the new frontier of automotive manufacturing and is poised to completely change the paradigm of automotive production.

If we imagine bringing a feature update to one of today's vehicles, it will most likely take anywhere from one to seven years from the idea to when that feature is actually perceptible in the production vehicle; this takes so long because the vehicles produced up to this point have not been designed with frequent updates in mind [**?**]. Traditionally focused on physical functionality, the automotive industry has evolved from early electronic features such as airbags, vehicle stabilisation and braking systems to modern driver assistance and even automated driving. The current shift towards a digital experience is possible thanks to vehicle design that includes software integration as a fundamental part. Software should no longer be seen as an accessory to the vehicle, but as an integral part of the vehicle itself.

The simultaneous efforts of major automotive companies such as Bosch, Renault and Stellantis, in collaboration with leading computer developers such as Arm, BlackBerry and AWS, have given rise to the Software Defined Vehicle concept, which they define as "any vehicle that manages its own operations, adds functionality and enables new features primarily or entirely through software" [**?**].

The Software Defined Vehicle solution is nowadays being considered by several companies as the manifesto of a new era of vehicle development. An example is given by the Renault Group, which in an overview of its products describes: "Today, it is already possible to make remote updates of some vehicles via the Firmware Over The Air (FOTA) system. This keeps the vehicle safe by making it

easier and faster to improve the on-board system and apply patches. Tomorrow, the Software Defined Vehicle's flexible and scalable architecture will enable the faster development and integration of new features throughout the vehicle lifecycle, directly into the cloud, that is, in secure online servers accessible from anywhere and anytime" [**?**].

It is evident that Software Defined Vehicles represent the future of the automotive industry, promising an enriched and sustainable user experience as vehicle technologies evolve. This section further clarifies the current state of the industry, highlighting the key enablers that are allowing the development of the SDV paradigm and the benefits of this innovation.

2.2.1 Enablers

There are mainly three fundamental technologies that contribute to the realisation of the Software Defined Vehicle: standardized hardware, cloud and over-the-air (OTA) updates via OTA servers, all developed by leading companies in the computing industry. In this section, each technology will be analysed with reference to concrete examples from the current market.

- **Standardized hardware**

  One of the most important aspects of Software Defined Vehicle is the separation of software from hardware. To achieve this, it is essential to move away from the approach of using dedicated hardware for each vehicle component system, and instead favour an approach based on general purpose processors that are as centralised as possible. This transition not only promotes ease of software development and scalability, but also offers the opportunity to create parity between the virtual development and test environment and the real execution environment.

  Several players in the semiconductor industry have stepped up to the challenge of realising this vision, including Arm. Through the development of energy-efficient processors, Arm is present in every part of the vehicle, from high-performance systems in advanced driver assistance systems (ADAS), automated driving (AD), in-vehicle infotainment (IVI) and digital cockpits, to gateway, body and microcontroller endpoints [**?**]. The aim is to create Arm-based MCUs that enable implementation of a common architecture, scalability between applications to meet processing requirements, software reuse and reduced development costs.

  Another major player is Qualcomm, which is being adopted by the Renault Group through its Snapdragon Digital Chassis vehicle architecture, a set of cloud-connected platforms for telematics and connectivity, digital cockpits, assistance and driver autonomy.

- **Cloud**

  Using a cloud platform that offers scalable and secure solutions for real-time application updates, increased connectivity and efficient data management is essential for SDV.

Well-known companies such as Amazon Web Services (AWS) and Google Cloud are already present in the automotive industry as partners of partner of many automotive companies. The AWS services and technologies will be in depth described in the futher chapters.

- **Over-The-Air updates**

  An Over-The-Air (OTA) update is the remote and wireless transfer of applications, services, firmware and configurations from a server to a target device. This process takes place over an available network, preferably the Internet. The main purposes of OTA are to remotely update software or firmware, provide power-safe procedures to ensure that the device will boot even if power is lost during the update process, maintain a robust implementation, ensure data protection and reduce overall maintenance costs [**?**].

  In the context of the thesis, it is crucial to acknowledge that the implementation of OTA updates may increase the vulnerability of automotive systems to hacking and other cyber attacks. These vulnerabilities could potentially be exploited by hackers to gain unauthorised access to private information, take remote control of the vehicle or even cause it to malfunction. Another significant issue is the leakage of information about updates and their sources. This can enable malicious actors to introduce viruses and malware, further exacerbating the security risks associated with OTA updates [**?**].

  To perform an OTA update, both a client on the vehicle, responsible for waiting and checking for incoming updates, and a server, facilitating the availability of the update broadcast to all connected devices, are essential. In this context, Autosar can be considered, as it represents a standard and open source architecture for intelligent mobility [**?**], which includes a dedicated platform for client and server management of OTA updates. Another notable example is Hawkbit, which serves as a backend framework for deploying software updates to edge devices and is being developed by the Eclipse Foundation; this tool will be discussed in more detail in later chapters as it will be used to create a proof of concept. The final tool of note is AWS Greengrass, an edge agent manager for managing software updates in edge IoT devices, provided by AWS; this tool will also be discussed in later chapters as an alternative solution to the client manager.

- **MQTT communication**

  The Message Queuing Telemetry Transport (MQTT) is a standardized protocol, specified by ISO/IEC 20922:2016 and developed by the Oaesis organization. It enables the exchange of Application Messages over a network connection, providing an ordered, lossless stream of bytes from the Client to Server and Server to Client without the need to support of a specific transport protocol.

  In an MQTT transport, an Application Message carries payload data, a Quality of Service (QoS), a collection of Properties, and a Topic Name. Clients, which can be programs or devices, perform various actions such as opening and closing network connections, publishing Application Messages, subscribing to requested Application Messages, and managing subscriptions [**?**].

On the Server side, it acts as an intermediary between publishing and subscribing Clients. The Server accepts network connections, processes Subscribe and Unsubscribe requests, and forwards Application Messages matching Client Subscriptions. The Server, also known as the Broker, essentially coordinates messages among various Clients. Its responsibilities extend to authorizing and authenticating MQTT Clients, transmitting messages to other systems for further analysis, and managing tasks such as handling missed messages and Client sessions [**?**].

Sessions, representing stateful interactions between Clients and Servers, can last for the duration of a Network Connection or span multiple consecutive connections.

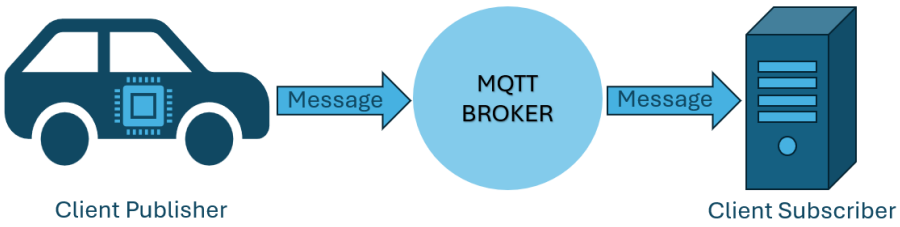

Figure 2.2. A simple representation of communication using the MQTT protocol

The MQTT protocol can be used in SDV, both for sending data produced by the vehicle to the cloud servers and for sending updates from the servers to the vehicle. This is because the MQTT protocol allows asynchronous and misaligned communication even in the presence of poor connectivity, a situation that cannot be underestimated in the automotive field.

The collaborative efforts of this technologies contribute to advancement of SDV for makeing vehicles not only defined by their physical attributes but also as dynamic entities that can be continuously updated through software.

2.2.2 Benefits

The Software Defined Vehicle, as introduced in the previous chapters, brings several benefits to both automotive companies and the end-user experience. These innovations are made possible by the fact that the vehicle becomes a device that can be constantly monitored and updated in real time via the cloud throughout its entire lifecycle. Let us now look at the key benefits.

From the point of view of this project, the main innovation brought by this technology is the security of the device software. Since, as mentioned above [**?**],

vehicles are considered as safety elements critical to human life, the safety benefits can be analysed from two perspectives:

- **Human Safety Critical Security:** The ability of SDV to receive real-time data from the vehicle allows in-depth monitoring of all its components. Taking the influence of tyres as an example, it has been found that most road accidents are caused by tyre wear and lack of regular maintenance. It is therefore necessary to assess the health of tyres through continuous monitoring of physical parameters such as tyre thickness, temperature and pressure, as well as regular maintenance. This helps to eliminate or minimise the possibility of tyre bursts and subsequent accidents. It also improves the safety of people and vehicles [**?**]. These factors can be monitored either manually or automatically: manual predictive maintenance requires human intervention and can lead to some errors; automatic predictive maintenance using artificial intelligence can be more efficient [**?**]. Renault defines this work as "predictive maintenance" [**?**], stressing the importance of collecting and analysing data in a centralised system to anticipate and prevent potential failures, ensure the safety of people, reduce maintenance costs and improve the performance of the vehicle.

- **Intrinsic Software Security:** In the presence of bugs and vulnerabilities in the vehicle's software, SDV makes it possible to intervene promptly to resolve each problem and reduce the window of exposure.

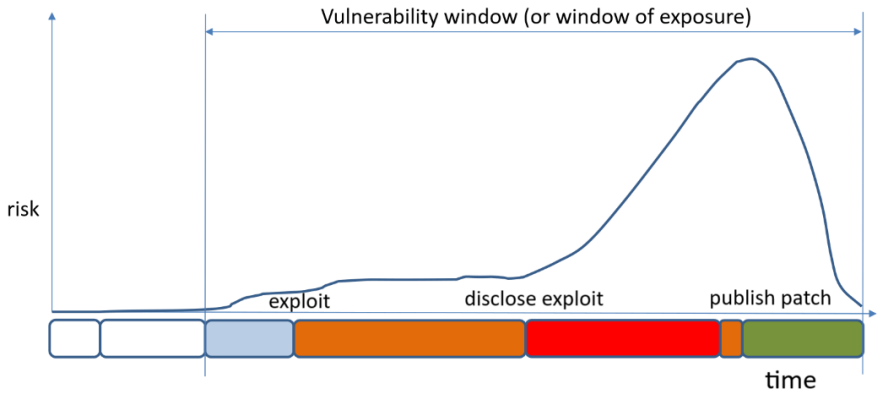

Figure 2.3. Risks and time relationship in the various phases of a vulnerability lifecycle

A crucial aspect of vehicle software security is the robustness of the algorithms, especially in the context of autonomous driving. In this context, a predictive algorithm responsible for vehicle safety decisions can be continuously improved and optimised. The SDV also introduces the concept of the 'digital twin', a platform that virtually replicates the functionality and behaviour of the vehicle. Thanks to this technology, predictive algorithms used in autonomous driving can be effectively tested on the cloud platform and, when ready, integrated directly into the vehicle.

From a user experience point of view, two other significant benefits can be identified: an increase in the value of the vehicle, which can be continuously upgraded over time, and the ability to enable additional vehicle functions via software. For example, the user can decide to activate a feature for a certain period of time and then deactivate it (paying only for the time it is used), or activate a new feature that was not available at the time of purchase. In essence, the vehicle becomes a dynamic platform that is constantly evolving and fully customisable through the software.

For automotive companies, the benefits mentioned so far can bring direct benefits to the industry. In support of this, Stellantis reports that: "the team in Poland will contribute to the global software creation network that is key to Stellantis' work in creating software-defined vehicles (SDVs) that offer customer-focused features throughout the vehicle's life span, including updates and features that will be added years after the vehicle is manufactured. "Creating an infrastructure inside our vehicles that easily and seamlessly adapts to meet driver expectations is a key element of Stellantis' global drive to deliver cutting edge mobility. Stellantis' software-driven strategy deploys next-generation tech platforms, building on existing connected vehicle capabilities to transform how customers interact with their vehicles and to generate €20 billion in incremental annual revenues by 2030".

In addition, the SDV paradigm brings an advantage from a software production pipeline perspective. In today's software production scenario, there can be two development mechanisms:

- A more traditional mode in which software is created directly on the system, hence on the processor itself. This is undoubtedly the most inconvenient solution, as it would require unnecessary overuse of processors, wasting resources, money and time.
- Alternatively, developers rely on cumbersome operating system emulation tools on the host machine and the cross-compilation process, which uses a dedicated compiler to produce executable code for the target system. Once the code is on the development system, a final integration and validation test can be performed, but scalability is limited to the number of physical hardware platforms.

Typical workflows for the development, integration and validation of embedded systems are as follows [**?**]:

As explained in the following chapters, by using the software defined vehicle, i.e. operating systems that rely on general purpose porpouse architectures to provide parity between cloud and edge systems, it is possible to reduce the embedded developer's workflow to remove many of the steps that are now no longer required, as shown in the diagram below [**?**]:

2.2.3 initiatives: SOAFEE

In 2021 Arm, AWS, and other founding members announced the Scalable Open Architecture for Embedded Edge (SOAFEE) Special Interest Group, which brings

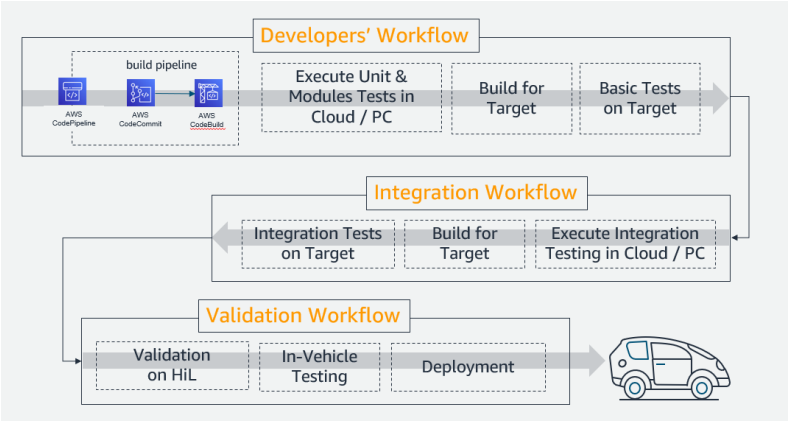

Figure 2.4. today development, integration, and validation workflows for embedded systems

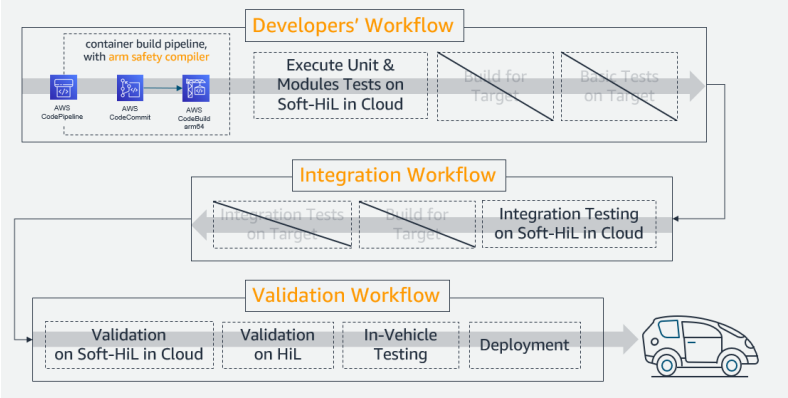

Figure 2.5. future development, integration, and validation workflows for embedded systems

together automakers, semiconductor, and cloud technology leaders to define a new open-standards based architecture to implement the lowest levels of a software-defined vehicle stack. [?]

SOAFE is created to achive Software Defined Vehicle, and for doing that four-pillar principle are used [?]:

1. Standards: standardization ensures interoperability and compatibility among various software components, fostering a cohesive ecosystem for Software Defined Vehicles.

2. New software architecture and methodologies: this involves transitioning from traditional monolithic architectures to more modular and scalable designs; the incorporation of agile development practices and DevOps methodologies ensures efficient and continuous software evolution.

3. Industry collaboration: Fostering partnerships, knowledge sharing and collaboration among key stakeholders, including automakers, technology companies and regulators, is essential.

4. Vehicle simulation: simulated environments allow in-depth testing and refinement of software functionality to ensure optimal performance and security under a variety of conditions.

SOAFEE aims to adopt and enhance current standards used in today's cloud-native world to help manage the software and hardware complexity of the automotive Software Defined Vehicle architecture.

The core principles of safety, security, and real time are inherent in each pillar. It is fully expected that the SOAFEE architecture will support use-cases that execute safety-critical services alongside non-safety-critical ones. It is fully expected that the SOAFEE architecture will support use cases that execute safety-critical services alongside non-safety-critical services. As it is not reasonable to develop the whole platform according to one safety standard, the strategy is to develop only safety-critical elements according to ISO 26262 and to isolate them from the non-safety-critical elements in order to ensure spatial, temporal and communication isolation. All implementations pass security checks and follow a set of best practices /citeSoafeeArchitectureOverview.

The SOAFEE paradigm is based on a very sophisticated architecture becouse it should work in the same way in the vehicle and in the cloud and follow cloud native technologies while considering the automotive specific needs for safety and limited resource footprints [**?**].

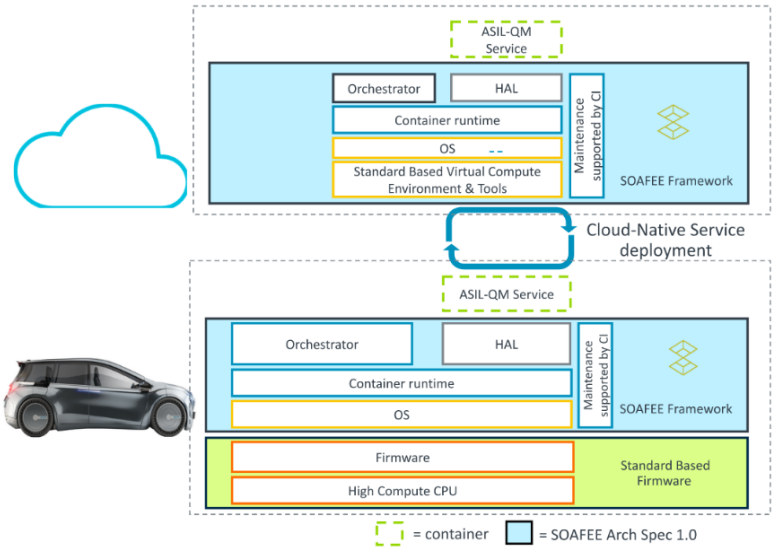

Figure 2.6. SOAFEE Architecture v1.0 [**?**]

Chapter 3

Cloud Computing and Amazon Web Services

As the analysis in previous chapters has shown, the Software Defined Vehicle is a pivotal advancement in the evolution of the entire automotive industry toward a safer, more efficient, and more sustainable future. Cloud computing is a crucial resource for SDV development due to its facilitation of development through its features and benefits. In the following section, cloud computing technologies will be analyzed in detail, focusing on one of the most important providers, Amazon Web Services.

3.1 Cloud Computing

The National Institute of Standards and Technology (NIST) provides the most comprehensive definition of cloud computing: "Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. This cloud model is composed of five essential characteristics, three service models, and four deployment models" [**?**]. This allows for a thorough analysis of the features of cloud computing in relation to AWS services, starting with the five essential characteristics.

- **On-demand self-service:** Consumers can access and allocate computing resources autonomously, such as server time and network storage, without direct involvement with service providers. AWS offers a vast cloud infrastructure with over 200 fully-featured services that consumers can easily access and use from their AWS account.

- **Broad network access:** Resources can be accessed over the network through standard mechanisms, making them usable across various client platforms. In AWS services, this is translated as an on-demand delivery of IT resources over the Internet with pay-as-you-go pricing.

- **Resource pooling:** Providers pool computing resources in a multi-tenant model, dynamically assigning them based on consumer demand. As said before AWS services are allocable e pagabili in base alle necessità del momento. The customer has limited control over the exact resource location but can specify a higher-level abstraction as country, state, or datacenter. In AWS, clients can select the geographic location of their services through regions. AWS Regions provide access to AWS services that are physically located in a specific geographic area. AWS provides the option to view the availability of a particular service in a specific region, in addition to selecting different regions [**?**]. Resources include storage, processing, memory, and network bandwidth. It also provides services for the Internet of Things, machine learning, data lakes, and analytics.

- **Rapid elasticity:** Resources can be easily adjusted to match fluctuations in demand, either automatically or manually. AWS provides various automated resource allocation systems, including the AWS Cloud Development Kit (AWS CDK) framework, which will be discussed later. The available capabilities are perceived as virtually limitless, and consumers can acquire them in any quantity at any time, always with a pay-per-use system.

- **Measured service:** Cloud systems efficiently manage resources through automated control and optimization, utilizing metering capabilities tailored to specific services such as storage, processing, bandwidth, and user accounts. For instance, AWS has infrastructure worldwide, allowing for easy deployment of applications in multiple physical locations. The proximity to end-users reduces latency and enhances their experience. This feature allows for clear and objective monitoring, control, and reporting of resource usage by both providers and consumers.

The three primary types of cloud computing are Infrastructure as a Service (IaaS), Platform as a Service (PaaS), and Software as a Service (SaaS). These options provide different levels of control, flexibility, and management, allowing users to configure services to meet their specific requirements.

- **Infrastructure as a Service (IaaS):** Consumers are able to utilize and deploy fundamental computing resources, including processing, storage, and networks. However, they only have control over operating systems, storage, and applications, as the cloud infrastructure is managed by the provider. Consumer control over some networking components is limited. Infrastructure as a Service (IaaS) provides a high level of flexibility and management control over IT resources. It is similar in practice to existing IT resources that many IT departments and developers are already familiar with.

- **Platform as a Service (PaaS):** Consumers can deploy their applications on the cloud infrastructure using the programming languages, libraries, services, and tools supported by the provider. The provider manages the underlying cloud infrastructure, including network, servers, operating systems, and storage, while consumers maintain control over their applications and configuration settings. This approach improves efficiency by eliminating the need

to manage resource procurement, capacity allocation, software maintenance, patching, or any other tasks involved in running your application.

- **Software as a Service (SaaS):** Consumers can use the provider's applications on the cloud infrastructure, which are accessible from different client devices through interfaces such as web browsers or programs. However, consumers do not have control over the underlying cloud infrastructure, including the network, servers, operating systems, and storage, except for limited user-specific application configuration settings. With a SaaS offering, users do not need to worry about maintaining the service or managing the underlying infrastructure. The focus should be on how to use the software effectively.

The analysis thus far has focused on cloud computing, specifically the essential characteristics that a cloud service must possess to be considered a true cloud service, as well as the service models that can be offered. Now, let's analyze in more detail the part related to cloud computing service deployment models and explore which models are most suitable for which workloads using AWS [**?**]. Note that in this case, there are slight differences between the NIST and AWS definitions of the various deployment modes.

- **public cloud:** According to NIST, a public cloud is defined as cloud infrastructure that is publicly accessible and owned, managed, and operated by businesses, academic institutions, government entities, or a combination thereof. In contrast, AWS defines a public cloud as infrastructure and services that are accessible over the public internet and hosted in a specific AWS Region.

- **private cloud:** Both NIST and AWS define private cloud as a cloud infrastructure exclusively provisioned for a single organization, which may own, manage, and operate it independently or in collaboration with a third party. However, there is a difference in the location of the infrastructure. According to NIST, the infrastructure can be located on or off premises, while in AWS documentation, the infrastructure is provisioned on premises using a virtualization layer.

- **hybrid cloud:** The hybrid cloud is a combination of two or more separate cloud infrastructures, private or public, connected by technology to facilitate data and application portability. It allows organizations to leverage the cloud for its efficiency and cost savings while also maintaining on-site security, privacy, and control.

Exploring the many benefits of cloud computing, the focus now shifts to a comprehensive analysis of the main value patterns of cloud computing, with some charts and graphs to help clarify the outlook.

When analyzing the resources required by a business, it becomes clear that satisfying demand through local services and resources often requires a monetary expenditure of resources that exceeds actual demand. This results in a wastage of resources that could be allocated more efficiently. On the other hand, the opposite

situation may also occur, as shown in **??** within the same domain, where there is a higher demand than the available resources. The adoption of cloud computing transforms this paradigm. With AWS, you can pay only for the computing resources you consume and exclusively for the amount you utilize. This approach enables more cost-effective and streamlined resource utilization.

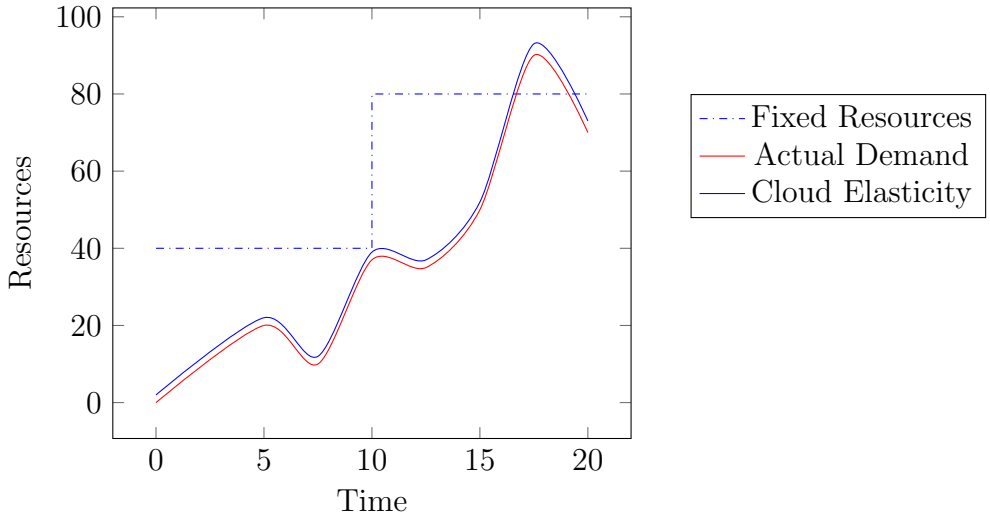

Table 3.1. Operating expenditure value model [**?**]

Cloud computing also reduces costs by aggregating resources needed by different companies in a transparent way to consumers. In addition to shortening the time to market and increasing earnings, cloud computing allows for access to resources anytime and anywhere, optimizing resource management with lower latency and a better experience as it is shown in refLocationFlexibilityValueModel.

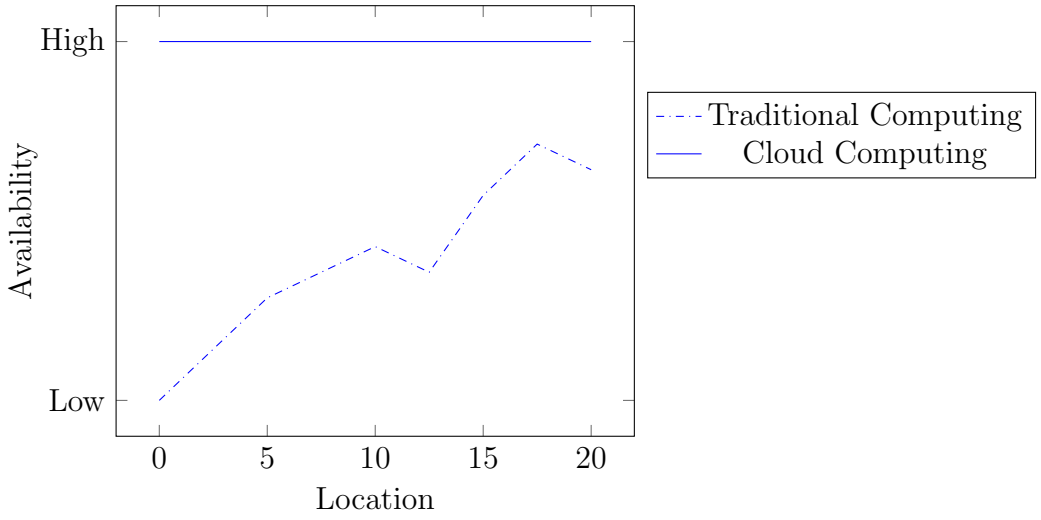

Table 3.2. Location flexibility value model [**?**]

In conclusion, cloud computing, exemplified by platforms like Amazon Web Services (AWS), enables organizations to access IT resources on-demand via the

Internet. This is facilitated by pay-as-you-go pricing models, which liberate organizations from the burdens of procuring, owning, and maintaining physical infrastructure. Cloud computing has a wide range of applications across industries, including the automotive sector. One of the main benefits of cloud computing is its dynamic scalability, which improves operational efficiency and reduces costs by utilizing resources more cost-effectively. This is due to the economies of scale inherent in cloud services, resulting in significantly lower variable expenses compared to self-managed infrastructure [**?**]. The characteristics of Amazon Web Services are analyzed in depth below.

3.2 Amazon Web Services

Amazon Web Services (AWS) is a widely adopted cloud solution with over 200 fully featured services available globally across multiple data centers. It is used by millions of customers, from emerging startups to industry giants and government agencies, as the cloud platform of choice to reduce costs, increase agility, and accelerate innovation [**?**].

AWS stands out by providing a broad set of services, including infrastructure technologies as well as cutting-edge capabilities such as machine learning, artificial intelligence, data lakes, analytics, and the Internet of Things. This extensive service portfolio facilitates the fast, easy and cost-effective migration of existing applications to the cloud and the creation of diverse digital solutions. AWS provides purpose-built databases for various application types, allowing users to choose the most suitable tool for optimal cost and performance. The depth of AWS services is unmatched, providing customers with a comprehensive toolkit for diverse computing needs.

Beyond its vast offerings, AWS has a large and dynamic global community with millions of active customers and tens of thousands of partners. This inclusive ecosystem spans industries and business sizes, with startups, enterprises, and public sector entities leveraging AWS for a myriad of use cases. The AWS Partner Network (APN) solidifies this network with thousands of system integrators and independent software vendors who adapt their technology to work on AWS.

AWS demonstrates its commitment to innovation through continuous technological advancements. In 2014, AWS launched AWS Lambda, which pioneered serverless computing. This allows developers to run their code without the need to provision or manage servers. Another example is Amazon SageMaker, a fully managed machine learning service that empowers developers to use machine learning without any previous experience.

Rooted in more than 17 years of operational experience, AWS offers unmatched reliability, security, and performance [**?**]. Since its establishment in 2006, AWS has become a globally trusted platform, revolutionizing IT infrastructure services by providing a highly reliable, scalable, and cost-effective cloud solution for businesses worldwide in the form of web services with pay-as-you-go pricing [**?**]. One of the main advantages of cloud computing is the ability to replace a company's initial capital expenditures required for infrastructure with low costs that vary as needed

and can scale with the business. AWS places great emphasis on the security of its systems and services, which is a fundamental pillar of their platform. In this thesis, we will analyze this feature in more detail.

3.2.1 Security

AWS is known for its flexible and secure cloud computing environment, designed to meet the strict security requirements of military, global banks, and high-sensitivity organizations.The infrastructure includes over 300 security, compliance, and governance services, supporting 143 security standards and compliance certifications. This architecture ensures scalability, reliability, and rapid deployment of applications and data while adhering to the highest security standards. Strong security at the core of an organization enables digital transformation and innovation. AWS utilizes redundant controls, continuous testing, and automation to maintain 24x7 monitoring and protection. Unlike customers' IT departments, which often operate on limited budgets, AWS prioritizes security as a core business aspect and allocates significant resources to safeguard the cloud and assist customers in ensuring robust cloud security. [**?**].

AWS empowers customers to confidently advance their businesses by providing a secure and innovative cloud infrastructure, a comprehensive suite of security services, and strategic partnerships. The AWS cloud infrastructure, combined with a comprehensive suite of security services and strategic partnerships, provides a solid foundation for secure innovation. Security is integrated and automated at every level of the organization, ensuring a swift and secure development process while reducing human errors.AWS offers a wide range of security services and partner solutions to help organizations effectively navigate evolving threats and compliance challenges. These expert-built capabilities equip organizations with the tools they need to stay secure and compliant [**?**].

The AWS global infrastructure follows rigorous security best practices and compliance standards, ensuring that users have access to one of the most secure computing environments in the world. It is designed and managed in alignment with a range of IT security standards, providing assurance to customers, including those in the life sciences industry, that their web architectures are built on exceptionally secure computing infrastructure. The main security standards obtained from infrastructure will now be explored through AWS documentation [**?**].

- **SOC 1, 2, 3:** AWS System and Organization Controls (SOC) Reports are third-party examination reports that demonstrate AWS's alignment with key compliance controls and objectives. SOC 1 focuses on controls relevant to a financial audit, covering security organization, access, data handling, change management, and more. SOC 2 expands to AICPA Trust Services Principles, evaluating controls related to security, availability, processing integrity, confidentiality, and privacy. The SOC 3 report is a publicly available summary of SOC 2. It includes an external auditor's assessment, AWS management's assertion, and an overview of AWS Infrastructure and Services. The report provides transparency and demonstrates AWS's commitment to security, compliance, and protection of customer data [**?**].

Figure 3.1. AWS System and Organization Controls Logo

- **FedRAMP:** FedRAMP is a US government program that ensures a standardized approach to security assessment, authorization, and continuous monitoring for cloud products and services. It is aligned with NIST SP 800 series. The program mandates that cloud service providers undergo an independent security assessment by a third-party assessment organization (3PAO) to verify compliance with the Federal Information Security Management Act (FISMA) [**?**].

- **ISO 9001:** "ISO 9001 is a globally recognized standard for quality management. It helps organizations of all sizes and sectors to improve their performance, meet customer expectations and demonstrate their commitment to quality. Its requirements define how to establish, implement, maintain, and continually improve a quality management system (QMS)" [**?**]. AWS ISO 9001:2015 certification directly supports customers developing, migrating, and operating their quality-controlled IT systems in the AWS cloud. They can use AWS compliance reports as evidence for their own ISO 9001:2015 programs and industry-specific quality programs [**?**].

- **ISO/IEC 27001:** ISO/IEC 27001 is a global security standard that outlines requirements for the systematic management of corporate and customer information. AWS has achieved ISO 27001 certification, demonstrating a comprehensive approach to assessing, managing, and mitigating information security risks. The certification covers AWS infrastructure, data centers, and services, ensuring ongoing compliance with international security standards [**?**].

- **ISO/IEC 27017:** "ISO/IEC 27017:2015 gives guidelines for information security controls applicable to the provision and use of cloud services by providing: additional implementation guidance for relevant controls specified in ISO/IEC 27002; additional controls with implementation guidance that specifically relate to cloud services. This Recommendation — International Standard provides controls and implementation guidance for both cloud service providers and cloud service customers" [**?**]. This certification ensures the

implementation of precise, cloud-specific controls and validates AWS commitment to robust security measures in cloud services [**?**].

- **ISO/IEC 27018:** ISO 27018 is a global code of practice for safeguarding personal data in the cloud. It builds upon ISO 27002 and offers guidance on implementing controls for Personally Identifiable Information (PII) in public clouds. AWS's ISO 27018 certification affirms its dedication to internationally recognized standards, emphasizing privacy and content protection [**?**].

- **HITRUST:** The Health Information Trust Alliance Common Security Framework (HITRUST CSF) integrates global standards such as GDPR, ISO, NIST, PCI, and HIPAA to establish a comprehensive framework for security and privacy controls. Some AWS services have been assessed under the HITRUST CSF Assurance Program by an approved HITRUST CSF Assessor and have been found to meet the HITRUST CSF Certification Criteria. Customers can inherit AWS certification for controls relevant to their cloud architectures established under the HITRUST Shared Responsibility Matrix (SRM). The certification is valid for two years, describes the AWS services that have been validated, and can be publicly accessed [**?**].

- **STAR:** The Cloud Security Alliance (CSA) introduced the Security, Trust, and Assurance Registry (STAR) to promote transparency in cloud provider security practices. STAR is a publicly accessible registry that documents the security controls of cloud computing offerings. AWS has joined the CSA STAR Self-Assessment, aligning with CSA best practices. The completed CSA Consensus Assessments Initiative Questionnaire (CAIQ) reports for AWS are publicly available [**?**].

Chapter 4

Proof Of Concept

This chapter explores the Proof of Concept (PoC) phase within the context of the thesis, aiming to validate the feasibility and efficacy of implementing SDV technologies. The main objective of this chapter is to translate theoretical concepts into concrete results, demonstrating the practical application of SDV in real-world situations. Through the PoC, we aim to confirm the fundamental principles and features of SDV, including its potential impact on vehicle performance, user experience, and overall safety.

The multifaceted nature of SDV requires a structured approach to its implementation, taking into account factors such as standardized hardware, cloud integration, and over-the-air (OTA) updates. To achieve this, a PoC was designed to address these components individually and holistically, ensuring a seamless integration that aligns with the envisioned paradigm shift in automotive manufacturing. Furthermore, this chapter aims to demonstrate the collaborative efforts with industry-leading technologies and platforms, highlighting the strategic partnerships forged with key players in the automotive and software development sectors. By aligning with renowned entities, the PoC aims to leverage their expertise, technologies, and frameworks, thereby enhancing the robustness and scalability of the SDV ecosystem.

Test and Validation are the concluding phases of this chapter, where the Proof of Concept is subjected to real-world scenarios. A demonstration involving a Raspberry Pi (RPi) serves as a tangible validation of the implemented SDV functionalities. This section serves as the litmus test, affirming the seamless orchestration of SDV within the envisioned architecture.

The exploration of the POC begins by detailing the services and technologies offered by Amazon Web Services (AWS) in the IoT, data management, and automotive essential for project implementation.

4.1 AWS Used Services

As discussed in previous chapters, the development of SDV technology requires a cloud infrastructure to handle server-side operations. AWS is a leading player in

the cloud world, and therefore an ideal alternative for the advancement of SDV, as well as an active partner in the implementation of technologies that contribute to the creation of a publicly available SDV for all. The following discussion introduces and analyzes, via AWS documentation the key tools for successful Proof of Concept (POC) implementation.

- AWS CLI: The AWS Command Line Interface (CLI) is an essential tool for developing with AWS services. t allows interaction with AWS services from the command line of a local PC, enabling the creation of infrastructure and management of properties from the command line.
- AWS Boto: Boto is an AWS SDK made for Python. A software Development Kit (SDK), more generally, is a set of creation tools specifically for developing and running software in a single platform. It includes resources such as documentation, examples, and APIs to facilitate faster application development. Boto basically works as an interface for applications that need to interact with and take advantage of the services provided by AWS. The AWS SDK for JavaScript v3 is another example of an SDK for JavaScript that works basically in the same way.

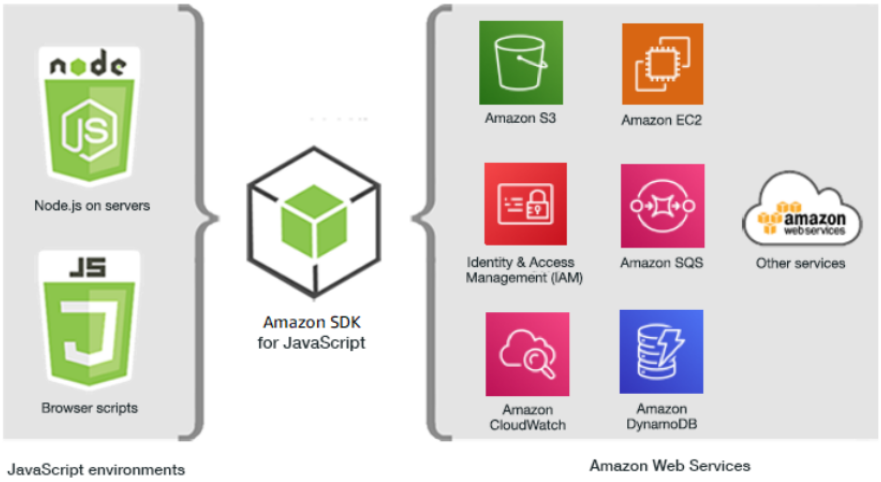

Figure 4.1. The high level rappresentation of the AWS SDK for JavaScript v3 [**?**]

- AWS CDK: The AWS Cloud Development Kit (CDK) "is an open-source software development framework for defining cloud infrastructure in code and provisioning it through AWS CloudFormation" [**?**]. This tool was used in the final phase of the POC design to automate the creation of the stack comprising all the services used.
- AWS IoT Core: AWS IoT Core provides the ability to connect IoT devices to AWS cloud services. AWS IoT Core enables the connection of IoT devices to AWS cloud services. It simplifies the integration of IoT devices with other AWS services. This is especially relevant in the automotive industry, where vehicle system ECUs can be viewed as multiple IoT devices. Communication between the device and AWS services can occur in several modes, with the

MQTT protocol being the most important for this project. The device can be connected by developing applications that utilize the SDK libraries. Once the data is transmitted, it can be utilized for various purposes such as testing, validation, and analysis.

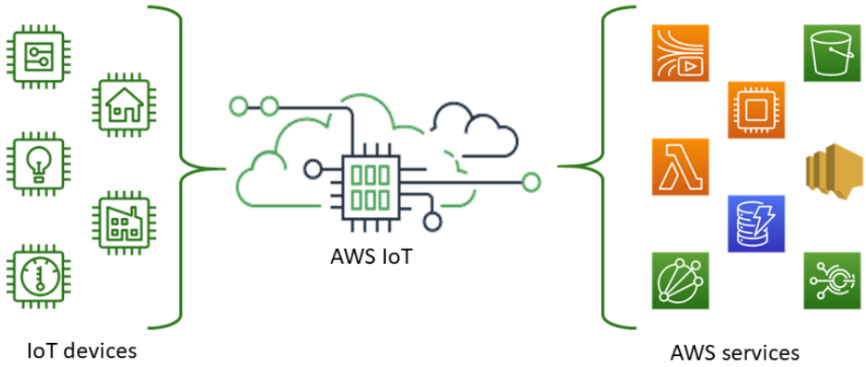

Figure 4.2. AWS IoT Core connection system beetween IoT device and AWS service [?]

- AWS IoT Greengrass: "AWS IoT Greengrass is an open source Internet of Things (IoT) edge runtime and cloud service that helps you build, deploy and manage IoT applications on devices" [?]. It is designed to work with intermittent connections and can manage fleets of devices in the field, locally or remotely, using MQTT or other protocols. Once installed, this service can be accessed through the command line. It was utilized in the early stages of project development as an agent to handle updates on the vehicle simulator side. However, this solution will be replaced by a custom solution as explained later.

- AWS IAM: "AWS Identity and Access Management (IAM) is a web service for securely controlling access to AWS services [...] such as access keys, and permissions that control which AWS resources users and applications can access" [?]. IAM is a service that provides a powerful access management mechanism. However, for the purpose of this thesis, only the relevant functionality to the project will be analyzed, specifically IAM's role management capabilities.An IAM role is an identity within AWS that can be assigned specific permissions via permission policies to determine what actions can and cannot be taken. Roles can be assumed by users, applications, or services that do not normally have access to the specific AWS resources. The IAM service also provides another important concept, that of policy, which is an AWS object that, when attached to an identity (including roles) or a resource, enables the creation of permissions and access control to other resources. For example, as explained below, a policy can be attached to the cloud representation of an IoT Core device to enable the connection of the physical dual IoT device or to grant Subscriber or Publisher permissions in a communication via MQTT protocol.

- AWS Lambda: AWS Lambda is a computing service that provides the ability

to run code without servers. It runs code on a high-performance computing infrastructure and handles administrative tasks related to computing resources autonomously, such as server and operating system management, capacity provisioning, automatic scaling, and logging. It is possible to run code for potentially any type of backend application or service [**?**]. Code can be written directly in Lambda console or imported from the local environment, and it supports several languages, including Python and JavaScript. The Lambda service can also manipulate data from other AWS services or manage tasks with services outside AWS as will be analyzed below.

- AWS Codepipeline: AWS CodePipeline is a fully managed continuous delivery service that automates release pipelines for software updates. It enables fast and reliable updates to applications and infrastructure, facilitating the rapid release of new features, iterative development based on feedback, and bug detection through testing every code change. The software release process can be modeled and configured quickly via the stages execution. A stage is a logical unit that creates an isolated environment and allows for the execution of a limited number of concurrent software changes. Each stage contains actions that are executed on application artifacts, such as source code from Codecommit. For instance, as shown in the image **??**, it is feasible to establish a software development pipeline that incorporates a codecommit repository as its source stage. This way, a codecommit-related event triggers the pipeline execution which then proceeds to the software build stage. An execution is defined as a series of modifications released from a pipeline. Each execution represents a set of modifications, such as a merged commit or a manual release of the last commit. Subsequently, the pipeline moves on to the test stage where the desired tests can be launched via Codebuild, and finally delivers the application for production.

Figure 4.3. An example of a CodePipeline in which some stages are reported [**?**]

- AWS Codebuild: AWS CodeBuild is a fully managed build service in the cloud that provides source code compilation, unit testing, and production of executable programs ready for distribution [**?**]. CodeBuild provides out-of-the-box configuration of compilation environments for popular programming languages, such as Python. It is also possible to create build platforms for programming languages for which there is no preconfiguration, but in this case it is necessary to leverage multiple AWS services. It is also possible to

use codebuild to run tests on application code using for example the pytest tool that allows you to test python code.

- AWS Codecommit: "AWS CodeCommit is a version control service that enables you to privately store and manage Git repositories in the AWS Cloud" [?]. This service becomes particularly interesting in the context of multiple services working together, including Lambda, Codepipeline, and Codebuild, because it allows the repository's Git and all its associated events (such as commit and push) to be used to trigger events that can automate various operations, such as triggering a pipeline in Codepipeline. As a result, CodePipelines typically use a CodeCommit repository as their inputo Source stage.

- Amazon S3: "Amazon Simple Storage Service (Amazon S3) is an object storage service that offers industry-leading scalability, data availability, security, and performance" [?]. The data saved in the storage is physically placed in multiple locations to ensure the durability of the data even if there is tampering with an item due to the presence of these copies; optionally, it can also be chosen to store the data in a single location to reduce the cost of the service. Amazon S3 can be used for data collection, aggregation, and analysis in many contexts and scenarios, but in the scope of this project, this service is used to store data that is transferred from one stage of the Codepipeline to another. Amazon's Codepipeline service automatically implements this method of output use. However, data stored in S3 from one stage to another can be manipulated through integration with other AWS services, such as Lambda.

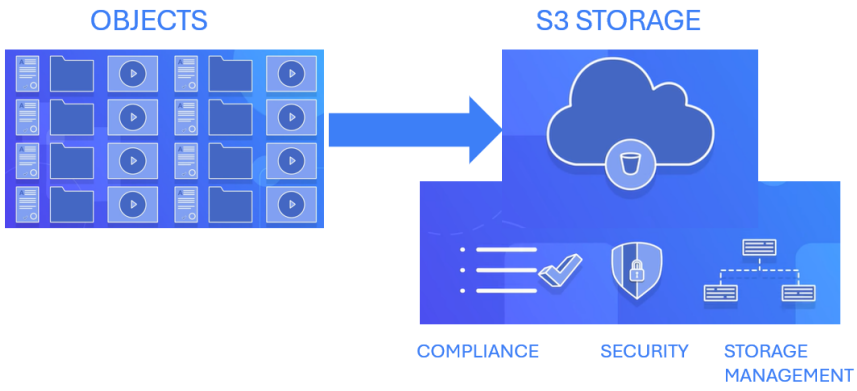

Figure 4.4. Amazon S3 high level storing rappresentation

- Amazon ECR: "Amazon Elastic Container Registry (Amazon ECR) is an AWS managed container image registry service that is secure, scalable, and reliable. Amazon ECR supports private repositories with resource-based permissions using AWS IAM. This is so that specified users or Amazon EC2 instances can access [...] container repositories and images" [?]. Basically, as shown in the figure ??, once the software has been produced and packaged, for example through the use of the CodeBuild service, it can be uploaded to Amazon ECR. The ECR takes care of encrypting the image and controlling

access to it, and then automatically manages the entire lifecycle of the image. Once the image is on ECR, it can be used either as an image for local download or through other AWS services.

Figure 4.5. Example of how Amazon ECR works in production and for pulling images [**?**]

- EC2: Amazon Elastic Compute Cloud (Amazon EC2) provides scalable, on-demand computing capacity in the Amazon Web Services (AWS) cloud. With Amazon EC2, users can create and use virtual machines in the cloud, instantiating resources as needed to perform compute operations. Amazon EC2 is a common choice for rapidly deploying applications because it provides an excellent computing resource at a low cost [**?**], and it is possible to manage networks of different instances of EC2 virtual machines through Amazon Virtual Private Cloud (VPC) and set their relative security, either on a per-instance basis or on an overall network basis. Additionally, it is possible to increase the capacity (scale up) of the instance after creation to handle computationally heavy tasks, such as spikes in website traffic. Conversely, if utilization decreases, capacity can be reduced (scale down). EC2 instances can be launched with Amazon Machine Images (AMIs), which are preconfigured templates containing the necessary components to use the server, including the operating system and additional software. AWS provides pre-built AMIs, but it is also possible to create your own AMIs using containers. Furthermore, it is possible to connect to an EC2 instance through various communication systems, such as using SSH keys provided at the time of instance creation.

- AWS System Manager: AWS Systems Manager is a service that provides visibility and control of the infrastructure on AWS. It allows users to view operational data from multiple AWS services and manage the automation of operational tasks across different AWS resources [**?**]. The AWS System Manager service is particularly relevant to the project due to its application management capability, namely the Parameter Store. Parameter Store is used to securely store configuration data and secrets, such as passwords, connection strings, and Amazon Machine Image (AMI) identifiers. Values are stored hierarchically by assigning hierarchical names to stored values using the "/" character, while maintaining the uniqueness of the name. For example, names such as Parameters/Parameter1, Parameters/Parameter2 can be used. In addition, it is possible to choose whether to store the data as plain text or encrypted data. Stored data can be retrieved directly from other services, for example, by interacting with Lambdas and SDK code functions.

- Amazon Kinesis Data Streams: Amazon Kinesis Data Streams is used to collect and process large streams of data records in real time, and eventually route them through other AWS services to various data collection and analysis applications, such as Amazon S3 as it is shown in the image **??**. "The delay between the time a record is put into the stream and the time it can be retrieved (put-to-get delay) is typically less than 1 second. In other words, a Kinesis Data Streams application can start consuming the data from the stream almost immediately after the data is added" [**?**].

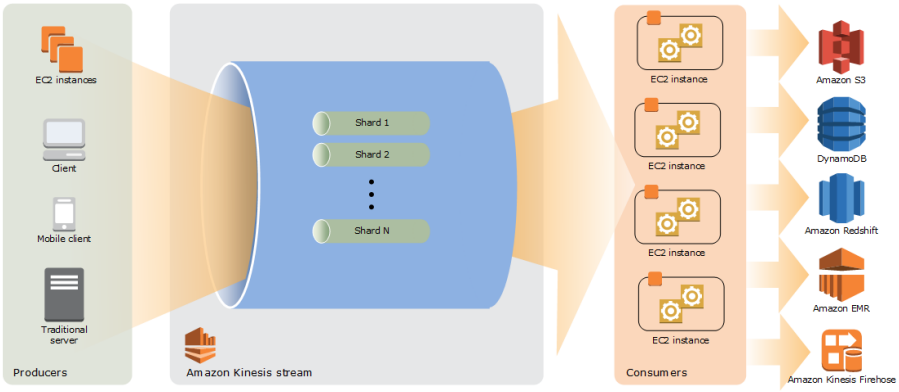

Figure 4.6. Illustration of the high-level architecture of Kinesis Data Streams with some examples of services that use the output of the stream. [**?**]

- Amazon Timestream: Amazon Timestream is a time-series database that allows you to store and easily analyze large amounts of data stored at regular intervals, ensuring that the time-series data is always encrypted, whether at rest or in transit. This service simplifies the complex process of managing the lifecycle of data by providing storage tiering with an in-memory store for current data and a magnetic store for historical data using Amazon S3 space. The transition of data between these two storage types is enabled through the use of user-configurable policies. The data lifecycle management mechanism makes Amazon Timestream ideal for handling telemetry data from IoT devices, for example. The service also provides a built-in interface for accessing data through a query engine [**?**]. The Timestrem service also provides an interface for Grafana to view and analyze stored data, which will be explored later.

- DinamoDB: Amazon DynamoDB is a full-featured NoSQL database service that provides high perfomances both speed and scalability. DynamoDB removes the administrative complexity of running and scaling your distributed database, so there's no need to manage provisioning, hardware setup and configuration, replication, software patching, or cluster sizing. DynamoDB also provides encryption at rest, eliminating the operational costs associated with protecting sensitive data. DynamoDB provides the ability to change the allocation of resources needed to store data in real time to use only the resources required. Additionally, DynamoDB offers on-demand backup functionality for long-term retention and archival purposes, as well as point-in-time recovery to safeguard against accidental write or delete operations. This feature

enables users to restore a table to any point within the last 35 days [**?**]. Note that this service was not utilized in the final version of the project, but was considered during development as an alternative for data storage and as a case study for understanding the data storage mechanisms used by AWS services.

- Amazon Cloudwatch: Amazon CloudWatch is a system to monitor the Amazon Web Services (AWS) resources and the applications running on the infrastructure in real time. With the use of CloudWatch it is possible to collect and track metrics from other AWS services such as Lambda, which are numeric variables that can be measured and analyzed for resources applications [**?**]. Practically this service represents the center for viewing and analyzing logs from the various AWS services in use.

All of the previously listed services have been useful, both as an active part in the project's realization and as potential options for the project's implementation, which will be analyzed below. The analysis of the POC will start with a description of the project's design, followed by an analysis of how the various services are actively involved in the project, and then an analysis of the interaction between the various elements that make up the architecture.

4.2 Architectural design

This section delves into the heart of the project implementation, starting with a high-level view of the various systems that make up its realization and transitioning to a visualization of the interactions between the various elements.

The study and analysis of the Software Defined Vehicle case study revealed that three fundamental elements were essential to create a concrete example of SDV implementation:

- TCU simulator: The TCU simulator is a system that replicates the basic functions of a telematic control unit (TCU). The system has the capability to send data packets and receive updates from a cloud server structure.
- Cloud infrastructure: The cloud infrastructure must be capable of managing both the data from the TCU and the update function.
- Data viewer: The data viewer is a platform that enables the visualization of manipulated and processed data in a manner that clearly displays changes in data behavior resulting from variations in the data and updates to the TCU.

With these three elements, a practical implementation made it possible to achieve what should happen in an SDV: having a vehicle system capable of updating itself via Over The Air updates. To support this implementation, it was necessary to use the services previously mentioned for creating the cloud infrastructure through the services made available by AWS.

Now, let's look at the components of the POC in detail through a high-level architectural representation of the previously introduced elements and a more detailed analysis of the code that compose them.

4.2.1 TCU Simulator

Come telematic control unit (TCU) in questo progetto si intende un sistema harware in grado di generare in qualche modo dei dati provenienti da una o più sottosistemi dei un eventuale veicolo, raccoglierli, per poi essere pronti per essere inviati all'esterno del veicolo. Per prendere familiarità con questo tipo di strutture e per dare forma al progetto di esempio, è stato necessario simulare una TCU in modo tale ceh fosse il più vicino possibile ad un sistema reale, senza però avere la disponibilità di un veicolo vero e proprio. Quindi riducendo ai minimi termini il concetto è stata presa come punto di arrivo di TCU una scheda raspberry pi che fosse in grado di generare dati telemetrici.

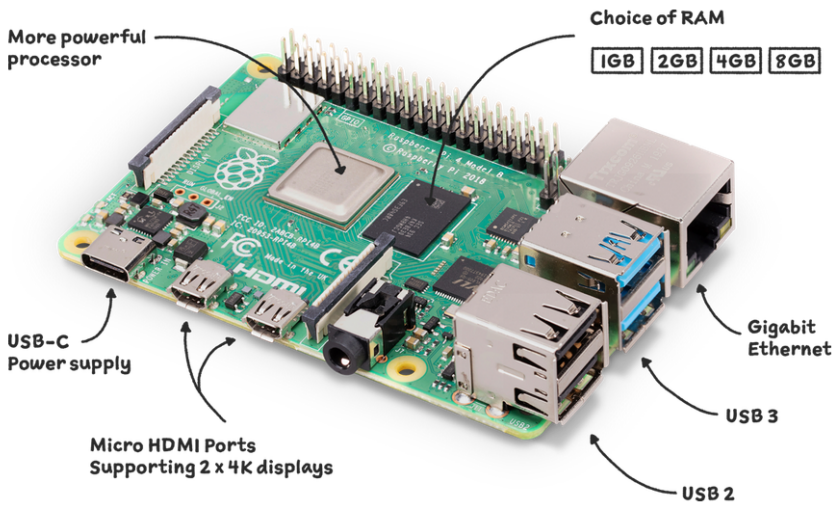

Figure 4.7. Illustration of a RaspberryPi board with its periferics [**?**]

Come riportato in figura **??** una RaspberryPi è una scheda contenente tutto il necessario per funzionare come un sistema indipendente a cui possono essere attaccate diverse periferiche, per questo motivo rappresenta l'ideale astrazione di una TCU general porpous necessaria per la realizzazione del caso di studio del SOC del SDV.

Prima di mettere mano alla raspberry però è stata necessaria una fase preliminare di realizzazione del progetto su macchina virtuale, il ciò ha facilitato notevolmente lo sviluppo e il test del codice scritto.

Il simulatore, nelle sue diverse fasi di sviluppo è stato pensato composto da 4 diverse componenti software, ognuna caratterizzata dalle proprie funzionalità:

- Elemento responsabile del collegamento al cloud server per l'invio di dati telematici;
- Elemento per il collegamento al server cloud per l'aggiornamento dell'unità telematica;
- Elemento per la generazione dei dati telematici;

- Elemento per l'update locale dell'unità telematica.

Questi 4 elementi sono orchestrati da uno script esterno che gestisce l'avvio dell'unità di simulazione tramite linea di comando. Analizzando nel dettaglio questi 4 elementi, il primo elemento in ordine cronologico preso come studio è stato l'elemnto responsabile al collegamento del sistema con l'infrastruttura cloud tramite protocllo MQTT. Questo elemento, ovviamente ha richiesto come azione preliminare la presenza di un servizio attivo di IoT Core lato cloud che come analizzeremo in seguito genera un certificato con annessa coppia di chiavi pubblica e privata per garantire l'identità del dispositivo. La disscussione relativa all'infrastruttura cloud verrà analizzata in seguito.

Per permettere il collegamento alla piattaforma cloud analizziamo il seguente codice **??** viene lanciato uno script Python che tramite la funzione AWSIoTMQTTClient della libreria AWSIoTPythonSDK.MQTTLib crea il collegamento, se possibile, con il server IoT Core per poi potergli inviare i dati. In particolare con mqttc = AWSIoTMQTTClient(VIN) viene creato il Clien MQTT, mentre con mqttc.configureEndpoint(ENDP 8883) is configured the host name and port number the client tries to connect to e con mqttc.configureCredentials(ROOT$_C A_F ILEPATH, PRIVATE_K EY_F ILEPATH, CERT_F I$

Listing 4.1. MQTT connection to the IoT Core AWS service

```
from AWSIoTPythonSDK.MQTTLib import AWSIoTMQTTClient

certificate_path="./Permanent/Certificates/"

def telemetry_handler(): #Function that manage the telemetry
    data sending to IoTCore via mqtt protocol
    global mqttc
    global connection_event
    #Thing connection
    VIN = "HawkbitDevice001" ##This is the Thing name
    ENDPOINT =
        "a1tbrylx7y3p3t-ats.iot.eu-west-1.amazonaws.com"
    CERT_FILEPATH = f"{certificate_path}{VIN}.cert.pem"
    PRIVATE_KEY_FILEPATH =
        f"{certificate_path}{VIN}.private.key"
    ROOT_CA_FILEPATH = f"{certificate_path}root-CA.crt"
    mqttc = AWSIoTMQTTClient(VIN)
    # Make sure you use the correct region!
    mqttc.configureEndpoint(ENDPOINT, 8883)
    mqttc.configureCredentials(ROOT_CA_FILEPATH,
        PRIVATE_KEY_FILEPATH, CERT_FILEPATH)
    # Connect to the gateway
    if mqttc.connect():
        print("Connected␣to␣IoT␣core.␣Now␣the␣device␣sends␣its␣
            telemetry␣every␣1␣seconds")
        connection_event.set() #Send connected signal to the
            main
        publish_topic = f"device/{VIN}/telemetry"
```

Il secondo elemento in analisi utilizzato per la realizzazione del simulatore è l'elemento che permette la connessione con il OTA Server per permettere il riconoscimento e il downolad degli aggiornamenti OTA. Questo è stato possibile tramite l'utilizzo come base di partenza del "Device Simulator" messo a disposizione in maniera aperta dal Hawkbit stesso. In questo caso si tratta di uno script Java che sfrutta l'interfaccia del Server Hawkbit pr collegarsi al server e rimanere in ascolto per attendere la presenza di aggiornamenti dedicarti al dispositivco stesso. Il collegamento viene fatto in automatico all'avvio del codice tramite le Simulation properties impostate, in particolare ai fini del progetto sperimentale il nome del device viene impostato in modo statico in SimulationProperties.java come private String name = "HawkbitDevice001"; mentre l'ip del server a cui connettersi viene preso come input in DeviceSimulator.java da System.setProperty("spring.rabbitmq.host", newHost); in cui newHost contiene l'informazione presa come input.

Listing 4.2. Hawbit Device simulator source code

```
private static long getOverallRead(final
   CloseableHttpResponse response, final MessageDigest md,
   final String url) throws IOException {
    long overallread = 0L;
    String[] urlParts = url.split("/");
    File downloadFolder = new File("./TCU/downloads");
    if (!downloadFolder.exists()) { //If "Download" folder
       doesn't exist
        boolean created = downloadFolder.mkdirs();
        if (!created) {
            System.err.println("Error␣in␣the␣directory␣
               creation!");
        }
    }
    File downloadFile = new
       File("./TCU/downloads/"+urlParts[10]);
    try (FileOutputStream outputStream = new
       FileOutputStream(downloadFile);
    final BufferedOutputStream bos = new
       BufferedOutputStream(new
       DigestOutputStream(outputStream, md))) {
        try (BufferedInputStream bis = new
           BufferedInputStream(response.getEntity().getContent()))
           {
            byte[] buffer = new byte[8192]; //byte dimension
               from createBuffer of ByteStream.class
            int bytesRead;
            while ((bytesRead = bis.read(buffer)) != -1) {
                bos.write(buffer, 0, bytesRead); //Here only
                   for the md hash correctness.
                overallread += bytesRead;
            }
        }
```

```
        }
        return overallread;
    }
```

In seguito, come mostrato nelle parti salienti del codice sorgente di DeviceSimulatorUpdater.java, nel momento in cui il dispositivo riceve un segnale di download, dopo aver eseguito diversi controlli di sicurezza, questo viene avviato e i dati ricevuti vengono inseriti nella cartella specifica "./TCU/downloads/" tramite l'utilizzo di uno stream che prende i dati in arrivo dal collegamento e li inserisce nella cartella selezionata, con il nome del file ricavato dall'url di download.

in un primo momento come un elemento che potesse interagire con l'utente: sostanzialmente il simulatore si occupava di fare 3 cose:

- Catturare gli input da utente;
- Raccogliere i dati ed inviarli al cloud;
- Riconoscere la presenza di aggiornamenti, quindi restartare il sistema con l'aggiornamento avvenuto.

SPIEGARE OGNI ELEMENTO CON CODICE ANNESSO

PARLARE DEL FATTO CHE L'IMPLEMENTAZIONE DEL CODICE è STATA FATTA IN TRE FASI: PRIMA CON UNINTERFACCIA E POCHI DATI GENERATI, POI SENZA INTERFACCIA E UNA ATTIVAZIONE DIRETTA DELL'INVIO DI DATI CON PYTHON POI CON C++ POI CON C. SPIEGARE OGNI SINGOLO ELEMENTO DELLA TCU FACENDO RIFERIMENTO ANCHE ALLA SCHEDA ED EVENTUALMENTE DARE LA COMPOSIZIONE DI UN GUADRO CHE PIANO PIANO SI COSTRUISCE PER CREARE L'INTERA TCU

4.2.2 Cloud Infrastructure

4.2.3 Data Viewer

4.3 Test and Validation

4.3.1 RPi demo

Chapter 5

Conclding Remarks

5.1 Contribution Recaps

5.1.1 Have we meet the PoC goals?

5.2 Future Works

5.2.1 Transform the poc in a product

5.2.2 Virtual workbenches

5.2.3 Manage additional Use Cases (ML, Cockpit Apps, remote ECU etc..)